

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv(r"C:\\Users\\Admin\\Downloads\\fiat500_VehicleSelection_Dataset
df
```

Out[2]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
0	1.0	lounge	51.0	882.0	25000.0	1.0	44.907242	8.6115598
1	2.0	pop	51.0	1186.0	32500.0	1.0	45.666359	12.241889
2	3.0	sport	74.0	4658.0	142228.0	1.0	45.503300	11.41784
3	4.0	lounge	51.0	2739.0	160000.0	1.0	40.633171	17.634609
4	5.0	pop	73.0	3074.0	106880.0	1.0	41.903221	12.495650
...
1544	NaN	NaN	NaN	NaN	NaN	NaN	NaN	leng
1545	NaN	NaN	NaN	NaN	NaN	NaN	NaN	conc
1546	NaN	NaN	NaN	NaN	NaN	NaN	NaN	Null valu
1547	NaN	NaN	NaN	NaN	NaN	NaN	NaN	fi
1548	NaN	NaN	NaN	NaN	NaN	NaN	NaN	sear

1549 rows × 11 columns

```
In [3]: df.head()
```

Out[3]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
0	1.0	lounge	51.0	882.0	25000.0	1.0	44.907242	8.6115598
1	2.0	pop	51.0	1186.0	32500.0	1.0	45.666359	12.2418895
2	3.0	sport	74.0	4658.0	142228.0	1.0	45.503300	11.41784
3	4.0	lounge	51.0	2739.0	160000.0	1.0	40.633171	17.63460922
4	5.0	pop	73.0	3074.0	106880.0	1.0	41.903221	12.49565029

In [4]: `df.describe()`

Out[4]:

	ID	engine_power	age_in_days	km	previous_owners	lat	U
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	

In [5]: `df.info()`


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1549 entries, 0 to 1548
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    1538 non-null   float64
1   model                 1538 non-null   object
2   engine_power          1538 non-null   float64
3   age_in_days           1538 non-null   float64
4   km                    1538 non-null   float64
5   previous_owners       1538 non-null   float64
6   lat                   1538 non-null   float64
7   lon                   1549 non-null   object
8   price                 1549 non-null   object
9   Unnamed: 9            0 non-null      float64
10  Unnamed: 10           1 non-null      object
dtypes: float64(7), object(4)
memory usage: 133.2+ KB
```

In [6]: `df.columns`

Out[6]: Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners', 'lat', 'lon', 'price', 'Unnamed: 9', 'Unnamed: 10'], dtype='object')


```
In [7]: sns.pairplot(df)
```

```
Out[7]: <seaborn.axisgrid.PairGrid at 0x1c73a71ff40>
```




```
In [8]: sns.distplot(df['ID'])
```

```
Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x1c73bdc4b80>
```



```
In [9]: sns.heatmap(df.corr())
```

```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x1c73d45f460>
```



```
In [10]: df1=df.fillna(value=1)
```

```
In [11]: x=df1[['age_in_days']]
y=df1['km']
```

```
In [12]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.4)
```

```
In [13]: from sklearn.linear_model import LinearRegression  
lr= LinearRegression()  
lr.fit(x_train,y_train)
```


Out[13]: LinearRegression()

```
In [14]: print(lr.intercept_)
```

10107.496918795696

```
In [15]: prediction= lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[15]: <matplotlib.collections.PathCollection at 0x1c73df4c130>



```
In [16]: print(lr.score(x_test,y_test))
```

0.6909047947917997

```
In [17]: print(lr.score(x_train,y_train))
```

0.7039740472013478

```
In [18]: from sklearn.linear_model import Ridge,Lasso
```

```
In [19]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

Out[19]: Ridge(alpha=10)

```
In [20]: rr.score(x_test,y_test)
```

Out[20]: 0.6909047949621456


```
In [27]: from sklearn import metrics
print("Mean Absolute Error",metrics.mean_absolute_error(y_test,prediction))
print("Mean Squared Error",metrics.mean_squared_error(y_test,prediction))
print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
Mean Absolute Error 15844.679730165151
Mean Squared Error 507836061.28438807
Root Mean Squared Error: 22535.218243549098
```

In []: