

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv(r"C:\Users\Admin\Downloads\18_world-data-2023 - 18_world-data-2023.csv")
df
```

Out[2]:

	Country	Density\n(P/Km2)	Abbreviation	Agricultural Land(%)	Land Area(Km2)	Armed Forces size	Birth Rate	Calling Code
0	Afghanistan	60	AF	58.10%	652,230	323,000	32.49	93.0
1	Albania	105	AL	43.10%	28,748	9,000	11.78	355.0
2	Algeria	18	DZ	17.40%	2,381,741	317,000	24.28	213.0
3	Andorra	164	AD	40.00%	468	NaN	7.20	376.0
4	Angola	26	AO	47.50%	1,246,700	117,000	40.73	244.0
...
190	Venezuela	32	VE	24.50%	912,050	343,000	17.88	58.0
191	Vietnam	314	VN	39.30%	331,210	522,000	16.75	84.0
192	Yemen	56	YE	44.60%	527,968	40,000	30.45	967.0
193	Zambia	25	ZM	32.10%	752,618	16,000	36.19	260.0
194	Zimbabwe	38	ZW	41.90%	390,757	51,000	30.68	263.0

195 rows × 35 columns

In [3]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 195 entries, 0 to 194
Data columns (total 35 columns):
#   Column                                          Non-Null Count  Dtype
---  -
0   Country                                         195 non-null    object
1   Density
(P/Km2)                                         195 non-null    object
2   Abbreviation                                   188 non-null    object
3   Agricultural Land( %)                        188 non-null    object
4   Land Area(Km2)                               194 non-null    object
5   Armed Forces size                             171 non-null    object
6   Birth Rate                                    189 non-null    float64
7   Calling Code                                  194 non-null    float64
8   Capital/Major City                           192 non-null    object
9   Co2-Emissions                                188 non-null    object
10  CPI                                             178 non-null    object
11  CPI Change (%)                               179 non-null    object
12  Currency-Code                                180 non-null    object
13  Fertility Rate                               188 non-null    float64
14  Forested Area (%)                            188 non-null    object
15  Gasoline Price                               175 non-null    object
16  GDP                                             193 non-null    object
17  Gross primary education enrollment (%)        188 non-null    object
18  Gross tertiary education enrollment (%)       183 non-null    object
19  Infant mortality                             189 non-null    float64
20  Largest city                                  189 non-null    object
21  Life expectancy                              187 non-null    float64
22  Maternal mortality ratio                     181 non-null    float64
23  Minimum wage                                 150 non-null    object
24  Official language                            194 non-null    object
25  Out of pocket health expenditure             188 non-null    object
26  Physicians per thousand                      188 non-null    float64
27  Population                                    194 non-null    object
28  Population: Labor force participation (%)     176 non-null    object
29  Tax revenue (%)                              169 non-null    object
30  Total tax rate                               183 non-null    object
31  Unemployment rate                            176 non-null    object
32  Urban_population                             190 non-null    object
33  Latitude                                      194 non-null    float64
34  Longitude                                    194 non-null    float64
dtypes: float64(9), object(26)
memory usage: 53.4+ KB
```

In [4]: `df.describe()`

Out[4]:

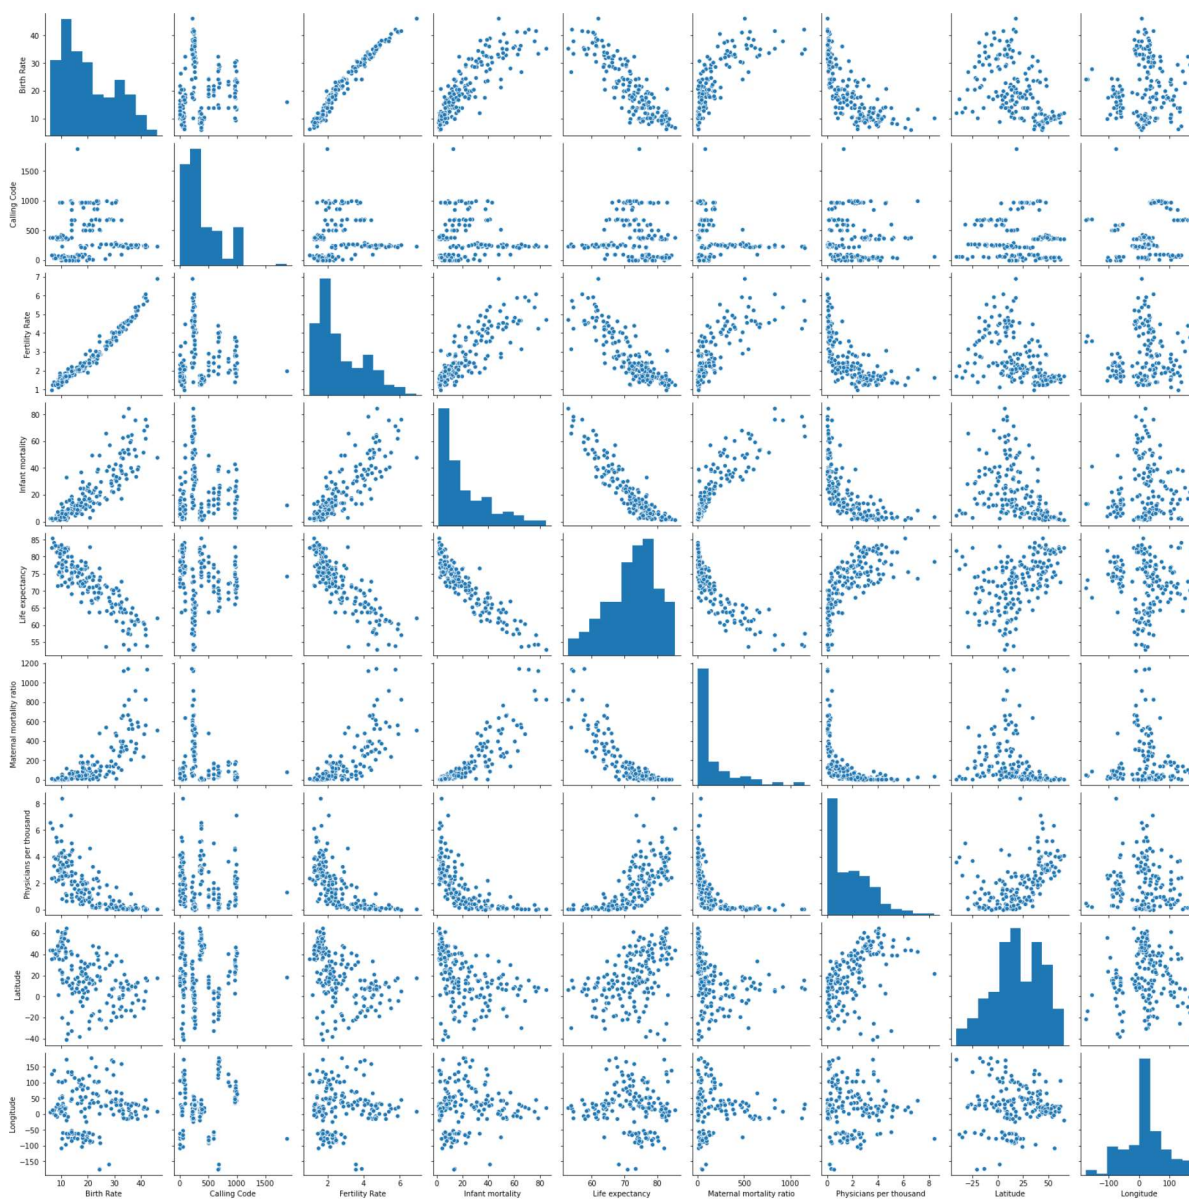
	Birth Rate	Calling Code	Fertility Rate	Infant mortality	Life expectancy	Maternal mortality ratio	Physicians per thousand	
count	189.000000	194.000000	188.000000	189.000000	187.000000	181.000000	188.000000	19
mean	20.214974	360.546392	2.698138	21.332804	72.279679	160.392265	1.839840	1
std	9.945774	323.236419	1.282267	19.548058	7.483661	233.502024	1.684261	2
min	5.900000	1.000000	0.980000	1.400000	52.800000	2.000000	0.010000	-4
25%	11.300000	82.500000	1.705000	6.000000	67.000000	13.000000	0.332500	
50%	17.950000	255.500000	2.245000	14.000000	73.200000	53.000000	1.460000	1
75%	28.750000	506.750000	3.597500	32.700000	77.500000	186.000000	2.935000	4
max	46.080000	1876.000000	6.910000	84.500000	85.400000	1150.000000	8.420000	6

In [5]: `df.columns`

Out[5]: Index(['Country', 'Density\n(P/Km2)', 'Abbreviation', 'Agricultural Land(%)', 'Land Area(Km2)', 'Armed Forces size', 'Birth Rate', 'Calling Code', 'Capital/Major City', 'Co2-Emissions', 'CPI', 'CPI Change (%)', 'Currency-Code', 'Fertility Rate', 'Forested Area (%)', 'Gasoline Price', 'GDP', 'Gross primary education enrollment (%)', 'Gross tertiary education enrollment (%)', 'Infant mortality', 'Largest city', 'Life expectancy', 'Maternal mortality ratio', 'Minimum wage', 'Official language', 'Out of pocket health expenditure', 'Physicians per thousand', 'Population', 'Population: Labor force participation (%)', 'Tax revenue (%)', 'Total tax rate', 'Unemployment rate', 'Urban_population', 'Latitude', 'Longitude'], dtype='object')

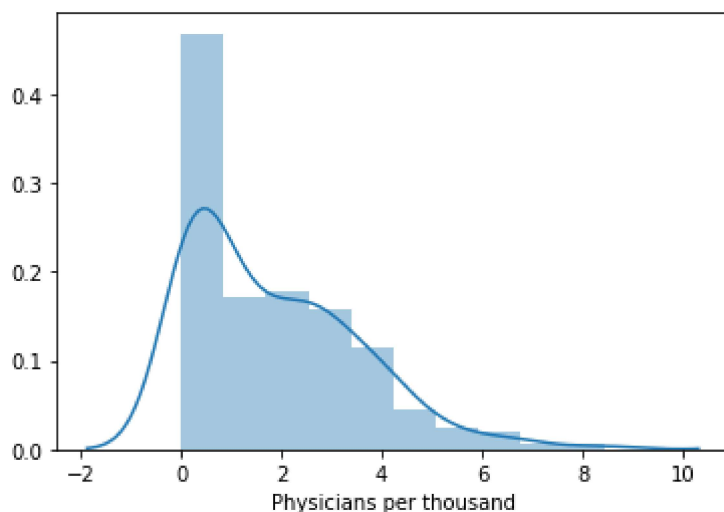
```
In [6]: sns.pairplot(df)
```

```
Out[6]: <seaborn.axisgrid.PairGrid at 0x216470c3df0>
```



In [7]: `sns.distplot(df['Physicians per thousand'])`

Out[7]: `<matplotlib.axes._subplots.AxesSubplot at 0x2164a4a93d0>`



In [8]: `df1=df[['Birth Rate','Calling Code','Fertility Rate','Infant mortality','Life expectancy','Maternal mortality ratio','Physicians per thousand','Latitude','Longitude']]`
`df1`

Out[8]:

	Birth Rate	Calling Code	Fertility Rate	Infant mortality	Life expectancy	Maternal mortality ratio	Physicians per thousand	Latitude	Longitude
0	32.49	93.0	4.47	47.9	64.5	638.0	0.28	33.939110	67.709953
1	11.78	355.0	1.62	7.8	78.5	15.0	1.20	41.153332	20.168331
2	24.28	213.0	3.02	20.1	76.7	112.0	1.72	28.033886	1.659626
3	7.20	376.0	1.27	2.7	NaN	NaN	3.33	42.506285	1.521801
4	40.73	244.0	5.52	51.6	60.8	241.0	0.21	-11.202692	17.873887
...
190	17.88	58.0	2.27	21.4	72.1	125.0	1.92	6.423750	-66.589730
191	16.75	84.0	2.05	16.5	75.3	43.0	0.82	14.058324	108.277199
192	30.45	967.0	3.79	42.9	66.1	164.0	0.31	15.552727	48.516388
193	36.19	260.0	4.63	40.4	63.5	213.0	1.19	-13.133897	27.849332
194	30.68	263.0	3.62	33.9	61.2	458.0	0.21	-19.015438	29.154857

195 rows × 9 columns

```
In [9]: sns.heatmap(df1.corr())
```

```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x2164af81c70>
```



```
In [10]: df2=df1.head(10)
```

```
In [11]: df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Birth Rate                           10 non-null     float64
1   Calling Code                         10 non-null     float64
2   Fertility Rate                       10 non-null     float64
3   Infant mortality                     10 non-null     float64
4   Life expectancy                      9 non-null      float64
5   Maternal mortality ratio             9 non-null      float64
6   Physicians per thousand              10 non-null     float64
7   Latitude                             10 non-null     float64
8   Longitude                           10 non-null     float64
dtypes: float64(9)
memory usage: 848.0 bytes
```

```
In [12]: x=df2[['Birth Rate','Calling Code','Fertility Rate','Infant mortality']]
         y=df2[['Physicians per thousand']]
```

```
In [13]: from sklearn.model_selection import train_test_split
         x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```


```
In [14]: from sklearn.linear_model import LinearRegression  
lr= LinearRegression()  
lr.fit(x_train,y_train)
```

Out[14]: LinearRegression()

```
In [15]: print(lr.intercept_)  
  
[7.51406493]
```

```
In [16]: prediction= lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[16]: <matplotlib.collections.PathCollection at 0x2164b3ab2e0>



```
In [17]: print(lr.score(x_test,y_test))  
  
-0.20174763826329856
```

```
In [18]: print(lr.score(x_train,y_train))  
  
0.8573500905913856
```

```
In [19]: from sklearn.linear_model import Ridge,Lasso
```

```
In [20]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

Out[20]: Ridge(alpha=10)

```
In [21]: rr.score(x_test,y_test)
```

Out[21]: 0.08119579878375216

```
In [22]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[22]: Lasso(alpha=10)

```
In [23]: la.score(x_test,y_test)
```

Out[23]: 0.3081204951917943

```
In [24]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[24]: ElasticNet()

```
In [25]: print(en.intercept_)
```

[5.56330459]

```
In [26]: print(en.predict(x_test))
```

[-1.07693353 2.5165552 2.90747588]

```
In [27]: print(en.score(x_test,y_test))
```

0.1535675887254958

Evaluation

```
In [29]: from sklearn import metrics
print("Mean Absolute Error",metrics.mean_absolute_error(y_test,prediction))
print("Mean Squared Error",metrics.mean_squared_error(y_test,prediction))
print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

Mean Absolute Error 1.9268086168097331
Mean Squared Error 3.8424144875769266
Root Mean Squared Error: 1.9602077664311317

```
In [ ]:
```